

# **Systems Programming with C**



**15-123**

**Systems Skills in C and Unix**

{ my \$x;

DIR  
├── 1  
├── 2  
├── 3  
└── 4  
;

\$x++;

\mkdir \$ARGV[0]/temp;

opendir (DIR, "\$ARGV[0]");  
foreach \$file (readdir (DIR)) {

if (\$file ne "temp") {

system("mv \$file \$file2")

&X (mv \$file \$file2); }

\mv \$ARGV[0]/\$file \$ARGV[0]/temp;

DIR  
├── temp  
├── 1  
├── 2  
└── 3

# Why Systems Programming?

- To access computers resources at a lower level using system calls

- Examples


- Managing files, processes, IPC etc..

- Managing Files

- In Unix, any I/O component is a file

- stdin, stdout, device files, sockets


- All files created, open, read the same way



scanf free  
read — low level


# What is a system call?

- A direct request to the operating system to do something on behalf of the program
- Typically programs are executed in user mode
- System call allows a switch from user mode to kernel mode



# Unix Kernel

- The core of the unix operating system
- Managing
  - Processes
  - Files
  - Networking etc..
- More details from OS courses



# in Kernel Mode

- All programs run in
  - user mode
    - can be replaced by another process at any time
  - kernel mode
    - cannot be arbitrarily replaced by another process.
- A process in kernel mode
  - can be suspended by an interrupt or exception.
- A C system call
  - A software instruction that generates an OS interrupt or **operating system trap**
  - Assembly instruction X080

# Using System Calls

- To manage
  - the file system
    - Open, creat, close, read
  - control processes
    - <sup>fork</sup>~~fork~~, exec
  - provide communication between multiple processes.
    - pipes

# **File Systems**



# Create System Call

RWX  
111 = 7  
000 = 0  
010 = 2  
001 = 1

*creat("guna", 0444);*

#include <fcntl.h>

int creat(char\* filename, mode\_t mode) *└ guna*

- The mode
  - is an octal number
    - **Example: 0444** indicates that r access for USER, GROUP and ALL for the file.
  - If the file exists, the creat is ignored and prior content and rights are maintained.

# Opening Files

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(char* filename, int flags, mode_t mode);
```

- Flags: O\_RDONLY, O\_WRONLY, O\_RDWR, O\_CREAT, O\_TRUNC, O\_APPEND
- Mode: Specifies permission bits of the file
  - S\_IRUSR, S\_IWUSR, S\_IXUSR – owner permission
  - S\_IRGRP, S\_IWGRP, S\_IXGRP – group permission
  - S\_IROTH, S\_IWOTH, S\_IXOTH – other permission

*Handwritten example:*  
`open("ghn2", O_RDONLY | O_WRONLY, 0444)`  
*(Note: The original image has some scribbles over the flags and mode in the handwritten example.)*

# More on open

- Each open call generates a file descriptor (by kernel)
- Kernel keeps track of all open files
  - Up to 16 in general
- Each unix shell starts with 3 standard files
  - stdin (descriptor 0)
  - stdout (descriptor 1)
  - stderr (descriptor 2)
- All other file descriptors are assigned sequentially

*int fd = open("gnu", - , -);*

# Reading/Writing Files

- Low level read and write

- #include <unistd.h>

- ssize\_t read(int fd, void \*buf, size\_t n);

– Returns num bytes read or -1

- ssize\_t write(int fd, const void \*buf, size\_t n);

– Returns num bytes written or -1

char buf[20];

read(0, buf, 19);  
↑     ↑     ↑  
stdin

write(1, buf, ~~19~~);

↳ loop printf("%c", buf[i]);

↳ loop scanf("%c", &buf[i]);

↓  
fprintf(stdin, "%c", &buf[i]),

# lseek function

- `#include <sys/types.h>` *lseek(3, 54, 0);*
- `#include <unistd.h>`
- lseek moves the cursor to a desired position

**long lseek(int fd, int offset, int origin)**

origin	position
0	beginning of the file
1	Current position
2	End of the file

End of the file



- **Examples**


*lseek(3, -3, 2);*

# Closing a file

- include <unistd.h>
- int close(int fd);
  - Return 0 (success)
  - Return -1 (error)

# Example

```
int main(void){  
    char c;  
    while (read(0,&c,1) != 0)  
        write(1, &c, 1);  
    exit(0);  
}
```



- What does it do?

# Example

```
int foo(char s[], int size){  
    char* tmp = s;  
    while (--size>0 && read(0,tmp,1)!=0 &&  
        *tmp++ != '\n');  
    *tmp = '\0';  
    return (tmp-s);  
}
```

- What does it do?



# What about `size_t` and `ssize_t`

- `size_t` – unsigned int
- `ssize_t` - signed int
- How does this affect the range of values in each type?
  - with 32-bit int?

# What can go wrong with read and write?

- processing fewer bytes than requested
  - reaching EOF
  - Reading text lines from stdin
  - Reading and writing network sockets
    - Network delays
    - Buffering constraints

# Reading file metadata

- How can we find information about a file
- **#include <unistd.h>**
- **#include <sys/stat.h>**
- **int stat(const char\* filename, struct stat \*buf);**
- **int fstat(int fd, struct stat \*buf);**

# What is struct stat?

```
struct stat
{
    dev_t      st_dev;      /* ID of device containing file */
    ino_t      st_ino;      /* inode number */
    mode_t     st_mode;     /* protection */
    nlink_t     st_nlink;   /* number of hard links */
    uid_t      st_uid;     /* user ID of owner */
    gid_t      st_gid;     /* group ID of owner */
    dev_t      st_rdev;     /* device ID (if special file) */
    off_t      st_size;     /* total size, in bytes */
    blksize_t   st_blksize; /* blocksize for filesystem I/O */
    blkcnt_t    st_blocks;  /* number of blocks allocated */
    time_t     st_atime;    /* time of last access */
    time_t     st_mtime;    /* time of last modification */
    time_t     st_ctime;    /* time of last status change */
};
```

# Accessing File Status

**stat(char\* file, struct stat \*buf);**

**fstat(int fd, struct stat \*buf);**

**struct stat buf; // defines a struct stat to hold file information**

**stat("filename", &buf) ; // now the file information is placed in the buf**

**st\_atime** --- Last access time

**st\_mtime** --- last modify time

**st\_ctime** --- Last status change time

**st\_size** --- total size of file

**st\_uid** – user ID of owner

**st\_mode** – file status (directory or not)

# Example

```
#include <sys/types.h>
#include <sys/stat.h>
#include <dirent.h>
struct stat statbuf;

char dirpath[256];
getcwd(dirpath,256);
DIR *dir = opendir(dirpath);
struct dirent *dp;

for (dp=readdir(dir); dp != NULL ; dp=readdir(dir)){
    stat(dp->d_name, &statbuf);
    printf("the file name is %s \n", dp->d_name);
    printf("dir = %d\n", S_ISDIR(statbuf.st_mode));
    printf("file size is %ld in bytes \n", statbuf.st_size);
    printf("last modified time is %ld in seconds \n", statbuf.st_mtime);
    printf("last access time is %ld in seconds \n", statbuf.st_atime);
    printf("The device containing the file is %d\n", statbuf.st_dev);
    printf("File serial number is %d\n\n", statbuf.st_ino);
}
```

# How to determine a file type

- S\_ISREG
  - A regular file?
- S\_ISDIR
  - Is a directory?
  - `printf("dir = %d\n", S_ISDIR(statbuf.st_mode));`
- S\_ISSOCK
  - A network socket

# Working Directory

```
#include <unistd.h>
```

```
char* getcwd(char * dirname, int );
```



# Accessing Directories

```
struct dirent *readdir(DIR* dp)
```

returns a pointer to the next entry in the directory. A NULL pointer is returned when the end of the directory is reached. The struct dirent has the following format.

```
struct dirent {  
    u_long d_ino;                /* i-node number for the dir  
        entry */  
    u_short d_reclen;          /* length of this record */  
    u_short d_namelen ;         /* length of the string in  
        d_name */  
    char d_name[MAXNAMLEN+1] ; /* directory name */  
};
```

# Creating and removing Directories

- **int mkdir(char\* name, int mode);**
- **int rmdir(char\* name);**
  - returns 0 or -1 for success or failure.
- **mkdir("newfiles", 0400);**
- **rmdir("newfiles");**

# Example

```
#include <string.h>
#include <sys/types.h>
#include <sys/dir.h>

int search (char* file, char* dir){
    DIR *dirptr=opendir(dir);
    struct dirent *entry = readdir(dirptr);
    while (entry != NULL) {
        if ( strlen(entry->d_name) == strlen(file) && (strcmp(entry->d_name, file) == 0)
            return 0; /* return success */
        entry = readdir(dirptr);
    }
    return 1; /* return failure */
}
```

# File Management summary

- **creat( ), open( ), close( )**
  - managing I/O channels
- **read( ), write( )**
  - handling input and output operations
- **lseek( )**
  - for random access of files
- **link( FILE1, FILE2), unlink( FILE)**
  - aliasing and removing files
- **stat( )**
  - getting file status
- **access( ), chmod( ), chown( )**
  - for access control
  - `int access(const char *pathname, int mode);`
- **chdir( )**
  - for changing working directory
- **mkdir( )**
  - for creating a directory

# Dealing with system call interfaces

- System calls interface often change
  - place system calls in subroutines so subroutines
- Error in System Calls
  - returns -1
  - store the error number in a variable called “**errno**” given in a header file called **/usr/include/errno.h**.
- Using perror
  - When a system call returns an error, the function **perror** can be used to print a diagnostic message. If we call **perror( )**, then it displays the argument string, a colon, and then the error message, as directed by “**errno**”, followed by a newline.

```
if (unlink("text.txt")==-1){  
    perror("");  
}
```

# **Process Control**

# Process Control

- **exec( ), fork( ), wait( ), exit( )**
  - for process control
- **getuid( )**
  - for process ownership
- **getpid( )**
  - for process ID
- **signal( ) , kill( ), alarm( )**
  - for process control

# Other system functions

- **mmap(), shmget(), mprotect(), mlock()**
  - manipulate low level memory attributes
- **time(), gettimer(), settimer(), settimeofday(), alarm()**
  - time management functions
- **pipe( )**
  - for creating inter-process communication



# Coding Examples